

[My Programs](#) [Collaboration and Competition](#) [Submit Project](#)

## Project: Collaboration and Competition

### Submission Results

Submission Date: March 18, 2020

✓ Submission Passed

[Download Submission](#)

### Feedback Details

[Specification Review](#) [Code Review](#)

#### Reviewer Note

## Congratulations

### Great submission!

You have implemented and trained the agent successfully. All functions were implemented correctly and the [MADDPG \(Multi-Agent Deep Deterministic Policy Gradient\)](#) algorithm seems to work quite well.

Following posts give an insight into some other reinforcement learning algorithms that can be used to solve the environment

[Proximal Policy Optimization by Open AI](#)[Introduction to Various Reinforcement Learning Algorithms Part II \(TRPO, PPO\)](#)

#### Training Code

✓ Saved Model Weights

✓ Training code

#### Reviewer Note

You implemented a Multi Agent Deep Deterministic Policy Gradients algorithm, a very effective reinforcement learning algorithm. Your code was functional, well documented, and organized for training the agent.

Suggested reading:

- [Google Python Style Guide](#)
- [Python Best Practices](#)

Pros of the implementation

- Implementing the MADDPG algorithm is a good choice as it is found to work very well with multiple agents continuous action space.
- Correct Implementation of the Actor and Critic networks.
- Good use of replay memory to store and recall experience tuples.
- Using the target networks for Actor and Critic networks is a good choice.
- Good choice to update the target network using soft updates and using tau for it.

The repository includes functional, well-documented, and organized code for training the agent.

✓ Framework

#### README

✓ README.md

✓ Instructions

✓ Getting Started

✓ Project Details

#### Reviewer Note

Good work with environment details. The README described all the project environment details.

- **Environment:** How is it like?
- **Agent** and its **actions:** When is it considered resolved?; What are the possible actions the agent can take?
- **State space:** Is it continuous or discrete?
- **Reward Function:** How is the agent rewarded?
- **Task:** What is its task?; Is the task episodic or not?

The README describes the the project environment details (i.e., the state and action spaces, and when the environment is considered solved).

Report

- ✓

Report

▼
- ✓

Learning Algorithm

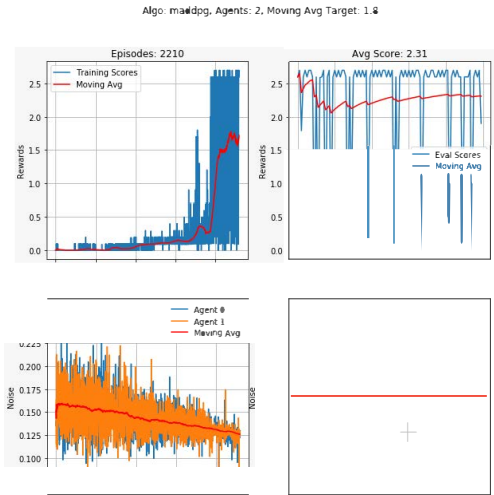
▼
- ✓

Plot of Rewards

▲

Reviewer Note

Nice work! The report informed the number of episodes needed to solve the environment and included a plot of rewards per episode.



A plot of rewards per episode is included to illustrate that the agents get an average score of +0.5 (over 100 consecutive episodes, after taking the maximum over both agents).

The submission reports the number of episodes needed to solve the environment.

- ✓

Ideas for Future Work